

```
In [ ]: from sklearn.preprocessing import scale
        from sklearn.datasets import load_iris
        from sklearn.cluster import KMeans
        import matplotlib.pyplot as plt
        import numpy as np
        import pandas as pd
        %matplotlib inline
```


```
In [ ]: from keras.datasets import mnist
        (train_x, train_y), (test_x, test_y) = mnist.load_data()
```

```
In [ ]: fig = plt.figure(figsize=(25, 4))
        for idx in np.arange(10):
            ax = fig.add_subplot(2, 10, idx+1, xticks=[], yticks=[])
            ax.imshow(train_x[idx], cmap='Blues_r')
            ax.set_title(str(train_y[idx]), fontsize=25)
```



```
In [ ]: img = train_x[0]

        fig = plt.figure(figsize = (12,12))
        ax = fig.add_subplot(111)
        ax.imshow(img, cmap='Blues_r')
        width, height = img.shape
        thresh = img.max()/2.5
        for x in range(width):
            for y in range(height):
                val = round(img[x][y],2) if img[x][y] !=0 else 0
                ax.annotate(str(val), xy=(y,x),
                            horizontalalignment='center',
                            verticalalignment='center',
                            color='white' if img[x][y]<thresh else 'black')
```



```
In [ ]: train_x = train_x.reshape(train_x.shape[0], -1)
print(train_x.shape)


sample_size = 5000
# Use only the top 1000 data for training
train_x = pd.DataFrame(train_x[:sample_size, :])
train_y = train_y[:sample_size]
```

(60000, 784)

```
In [ ]: from sklearn.decomposition import TruncatedSVD

tsvd = TruncatedSVD()
x_tsvd = tsvd.fit_transform(train_x)
markers=['o','v','^','<','>','8','s','P','*','X']
# plot in 2D by class
plt.figure(figsize=(10,10))
for i,marker in enumerate(markers):
    mask = train_y == i
    plt.scatter(x_tsvd[mask, 0], x_tsvd[mask, 1], label=i, s=10, alpha=1,marker=marker)
plt.legend(bbox_to_anchor=(1.00, 1), loc='upper left',fontsize=15)
```


Out[ ]: <matplotlib.legend.Legend at 0x261b799a340>



```
In [ ]: from sklearn.decomposition import NMF

nmf = NMF(n_components=2, init='random', random_state=0)
x_nmf = nmf.fit_transform(train_x)
markers=['o','v','^','<','>','8','s','P','*','X']
# plot in 2D by class
plt.figure(figsize=(10,10))
for i,marker in enumerate(markers):
    mask = train_y == i
    plt.scatter(x_nmf[mask, 0], x_nmf[mask, 1], label=i, s=10, alpha=1,marker=marker)
plt.legend(bbox_to_anchor=(1.00, 1), loc='upper left',fontSize=15)
```

Out[ ]: <matplotlib.legend.Legend at 0x261b78c20a0>




```
In [ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA(n_components=2)
x_lda = lda.fit_transform(train_x, train_y)

plt.figure(figsize=(10,10))
for i,marker in enumerate(markers):
    mask = train_y == i
    plt.scatter(x_lda[mask, 0], x_lda[mask, 1], label=i, s=10, alpha=1,marker=marker)
plt.legend(bbox_to_anchor=(1.00, 1), loc='upper left',fontSize=15)
```

Out[ ]: <matplotlib.legend.Legend at 0x2619c4e8580>




```
In [ ]: from sklearn.manifold import TSNE

tsne = TSNE(n_components=2)
x_tsne = tsne.fit_transform(train_x)

plt.figure(figsize=(10,10))
for i,marker in enumerate(markers):
    mask = train_y == i
    plt.scatter(x_tsne[mask, 0], x_tsne[mask, 1], label=i, s=10, alpha=1,marker=marker)
plt.legend(bbox_to_anchor=(1.00, 1), loc='upper left',fontsize=15)
```

Out[ ]: <matplotlib.legend.Legend at 0x261b5d9e6d0>




```
In [ ]: import umap

um = umap.UMAP()
x_umap = um.fit_transform(train_x)

plt.figure(figsize=(10,10))
for i,marker in enumerate(markers):
    mask = train_y == i
    plt.scatter(x_umap[mask, 0], x_umap[mask, 1], label=i, s=10, alpha=1,marker=marke
plt.legend(bbox_to_anchor=(1.00, 1), loc='upper left',fontsize=15)
```

```
c:\WProgramData\WAnaconda3\Wenvs\Wyujinyeo\lib\site-packages\numba\core\cpu.py:97: UserWarning: Numba extension module 'numba_dppy.numpy_usm_shared' failed to load due to 'ImportError(cannot import name 'TargetConfig' from 'numba.core.dispatcher' (c:\WProgramData\WAnaconda3\Wenvs\Wyujinyeo\lib\site-packages\numba\core\dispatcher.py))'.
numba.core.entrypoints.init_all()
```

```
Out[ ]: <matplotlib.legend.Legend at 0x261dd3ba850>
```



In [ ]:

In [ ]:

```
train = pd.read_csv(r"C:\temp\train.csv")

target = train['target']
train = train.drop(["target", "ID"], axis=1)

print ("Rows: " + str(train.shape[0]) + ", Columns: " + str(train.shape[1]))
train.head()
```

Rows: 4459, Columns: 4991

Out [ ]:

|   | 48df886f9 | 0deb4b6a8 | 34b15f335 | a8cb14b00 | 2f0771a37 | 30347e683 | d08d1fbe3 | 6ee66e115 |
|---|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 0 | 0.0       | 0         | 0.0       | 0         | 0         | 0         | 0         | 0         |
| 1 | 0.0       | 0         | 0.0       | 0         | 0         | 0         | 0         | 0         |
| 2 | 0.0       | 0         | 0.0       | 0         | 0         | 0         | 0         | 0         |
| 3 | 0.0       | 0         | 0.0       | 0         | 0         | 0         | 0         | 0         |
| 4 | 0.0       | 0         | 0.0       | 0         | 0         | 0         | 0         | 0         |

5 rows × 4991 columns

In [ ]:

```
from sklearn.preprocessing import StandardScaler
```

```
standardized_train = StandardScaler().fit_transform(train.values)
```

```
In [ ]: import chart_studio.plotly as py
import cufflinks as cf

import plotly.graph_objs as go
from plotly.offline import iplot

cf.go_offline()
cf.set_config_file(offline=False, world_readable=True)


In [ ]: feature_df = train.describe().T
feature_df = feature_df.reset_index().rename(columns = {'index' : 'columns'})
feature_df['distinct_vals'] = feature_df['columns'].apply(lambda x : len(train[x].va
feature_df['column_var'] = feature_df['columns'].apply(lambda x : np.var(train[x]))
feature_df['column_std'] = feature_df['columns'].apply(lambda x : np.std(train[x]))
feature_df['column_mean'] = feature_df['columns'].apply(lambda x : np.mean(train[x]))
feature_df['target_corr'] = feature_df['columns'].apply(lambda x : np.corrcoef(target
feature_df.head()
```

```
Out [ ]:
```

|   | columns   | count  | mean         | std          | min | 25% | 50% | 75% | max         | distinct_vals |
|---|-----------|--------|--------------|--------------|-----|-----|-----|-----|-------------|---------------|
| 0 | 48df886f9 | 4459.0 | 14654.930101 | 3.893298e+05 | 0.0 | 0.0 | 0.0 | 0.0 | 20000000.0  | 32            |
| 1 | 0deb4b6a8 | 4459.0 | 1390.894819  | 6.428302e+04 | 0.0 | 0.0 | 0.0 | 0.0 | 4000000.0   | 5             |
| 2 | 34b15f335 | 4459.0 | 26722.450922 | 5.699652e+05 | 0.0 | 0.0 | 0.0 | 0.0 | 20000000.0  | 29            |
| 3 | a8cb14b00 | 4459.0 | 4530.163714  | 2.359124e+05 | 0.0 | 0.0 | 0.0 | 0.0 | 14800000.0  | 3             |
| 4 | 2f0771a37 | 4459.0 | 26409.957390 | 1.514730e+06 | 0.0 | 0.0 | 0.0 | 0.0 | 100000000.0 | 6             |

```
In [ ]: feature_df = feature_df.sort_values('column_var', ascending = True)
feature_df['column_var'] = (feature_df['column_var'] - feature_df['column_var'].min())
trace1 = go.Scatter(x=feature_df['columns'], y=feature_df['column_var'], opacity=0.75)
layout = dict(height=400, title='Feature Variance', legend=dict(orientation="h"));
fig = go.Figure(data=[trace1], layout=layout);
iplot(fig);
```

## Feature Variance






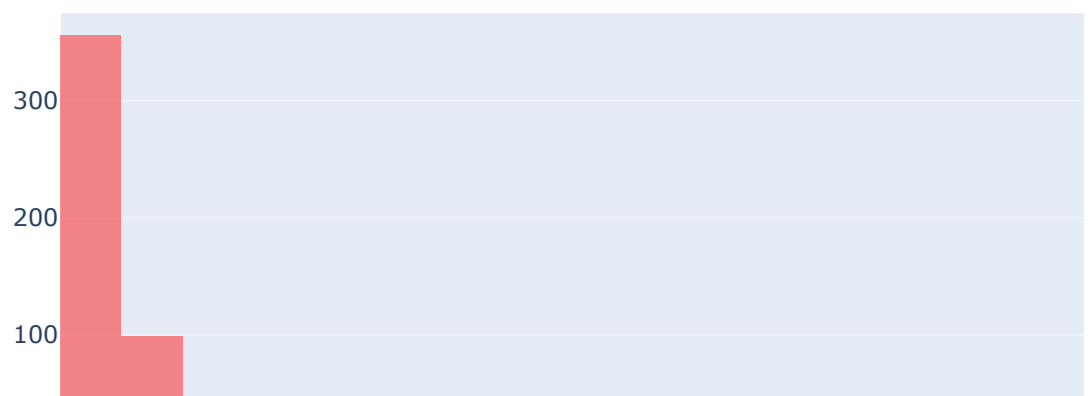
```
In [ ]: trace1 = go.Histogram(x=feature_df[feature_df['column_var'] <= 0.01]['column_var'], o
layout = dict(height=400, title='Distribution of Variable Variance <= 0.01', legend=d
fig = go.Figure(data=[trace1], layout=layout);
iplot(fig);

trace1 = go.Histogram(x=feature_df[feature_df['column_var'] > 0.01]['column_var'], op
layout = dict(height=400, title='Distribution of Variable Variance > 0.01', legend=di
fig = go.Figure(data=[trace1], layout=layout);
iplot(fig);
```

Distribution of Variable Variance <= 0.01




Distribution of Variable Variance > 0.01



```
In [ ]: trace1 = go.Histogram(x=feature_df['target_corr'], opacity=0.45, marker=dict(color="g",
layout = dict(height=400, title='Distribution of correlation with target', legend=dict(
fig = go.Figure(data=[trace1], layout=layout);
iplot(fig);
```

## Distribution of correlation with target



```
In [ ]: # Calculating Eigenvectors and eigenvalues of Cov matrix
mean_vec = np.mean(standardized_train, axis=0)
cov_matrix = np.cov(standardized_train.T)
eig_vals, eig_vecs = np.linalg.eig(cov_matrix)

# Create a list of (eigenvalue, eigenvector) tuples
eig_pairs = [(np.abs(eig_vals[i]), eig_vecs[:, i]) for i in range(len(eig_vals))]

# Sort the eigenvalue, eigenvector pair from high to low
eig_pairs.sort(key = lambda x: x[0], reverse= True)

# Calculation of Explained Variance from the eigenvalues
tot = sum(eig_vals)

# Individual explained variance
var_exp = [(i/tot)*100 for i in sorted(eig_vals, reverse=True)]
var_exp_real = [v.real for v in var_exp]

# Cumulative explained variance
cum_var_exp = np.cumsum(var_exp)
cum_exp_real = [v.real for v in cum_var_exp]

## plot the variance and cumulative variance
trace1 = go.Scatter(x=train.columns, y=var_exp_real, name="Individual Variance", opac
trace2 = go.Scatter(x=train.columns, y=cum_exp_real, name="Cumulative Variance", opac
layout = dict(height=400, title='Variance Explained by Variables', legend=dict(orient
```

```
fig = go.Figure(data=[trace1, trace2], layout=layout);  
iplot(fig);
```

## Variance Explained by Variables

